

# Distributed Systems & Network Principles

## Concurrency

Dr Vincent Gramoli | Lecturer  
School of Information Technologies



THE UNIVERSITY OF  
SYDNEY



- › Former lecture: a distributed system gives the illusion of a single system to a user thanks to transparency
- › Today's lecture: a single system gives the illusion of multiple resources to multiple users thanks to concurrency



- › A Very Brief History
- › UNIX Processes and Threads
- › Java Threads
- › Multithreading
- › Multithreaded Server
- › Virtualization



THE UNIVERSITY OF  
SYDNEY

# A Very Brief History

## A brief history

- › The hardware evolved
  - Hardware was expensive, humans were cheap (e.g., Multics 1960's)
  - Hardware cheaper, humans expensive (e.g., desktop computer 1980's)
  - Hardware very cheap, humans very expensive (e.g., handheld devices 2000's)
- › The OS interaction had to adapt
  - Batch: one execution at a time
  - Multiprogramming: multiple program executions simultaneously
  - Timeshare: split time into slots allocated for different program execution
  - GUI: multiple interfaces to access a system from different end-points
  - Ubiquitous devices: each user possess its computational device

# Uniprogramming vs. Multiprogramming

## Operating systems

### › Uniprogramming: one program execution at a time

- MS/DOS
- early Macintosh
- Batch processing
- No longer acceptable!

### › Many program executions on personal computers:

- Browsing the web
- Sending emails
- Typing a letter
- Burning a DVD

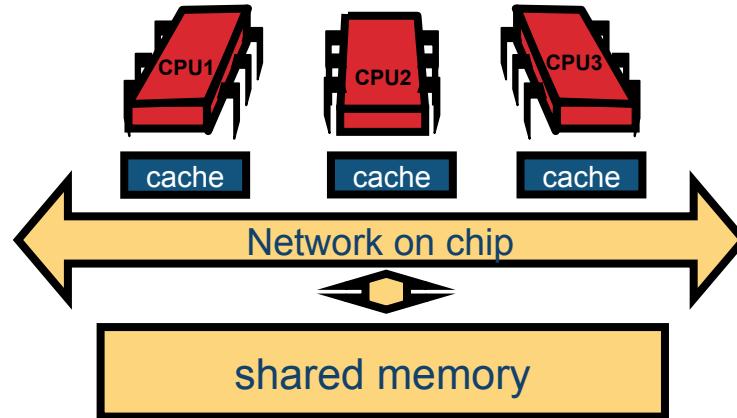
# Uniprogramming vs. Multiprogramming

## Operating systems

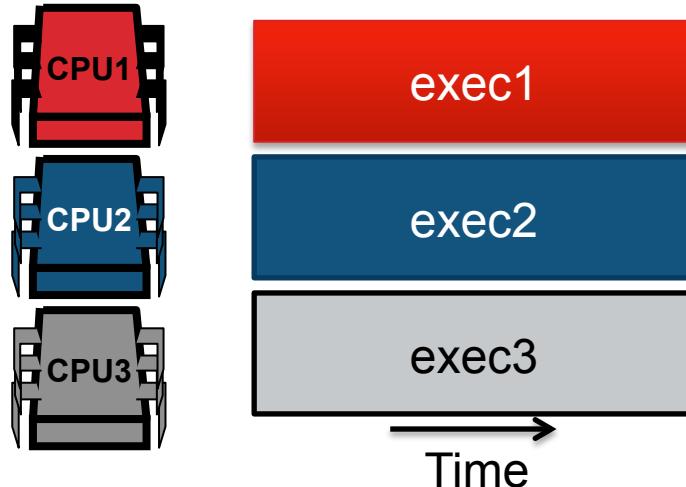
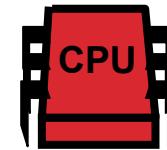
› Multiprogramming: more than one program execution at a time

- Multics
- UNIX/Linux
- Windows NT
- Mac OS X

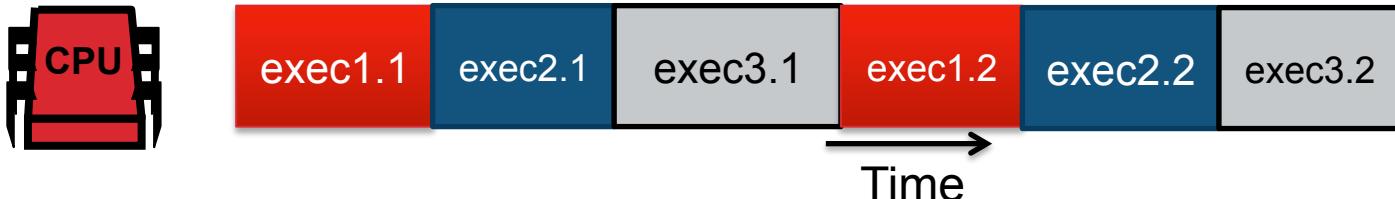
› Multi/many-cores



- › Assume a single central processing unit (CPU)
- › Problem: How to give the illusion to the users of multiple CPUs?



- › Solution: **scheduling** program executions, one after the other during small fractions of the time





THE UNIVERSITY OF  
SYDNEY

# Processes and Threads

Process: an abstraction representing a program execution

exec1

› When to switch from one process to another?

- timer interrupt goes off, explicit yield, I/O, etc.

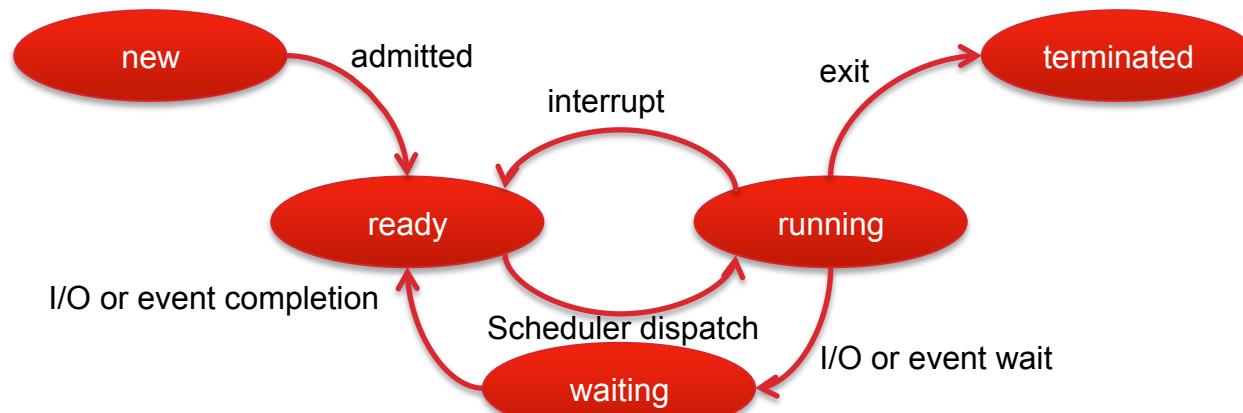
exec1.1 exec2.1

› Each virtual CPU needs to store the *process context*:

- program counter, stack pointer, other registers

Time →

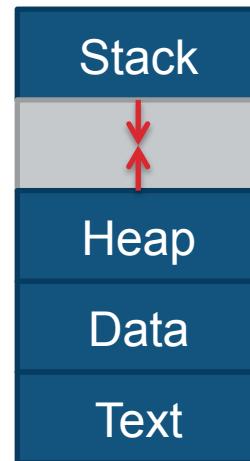
› A process switches from a state to another:



*Address space: A unit of management of a process's virtual memory*

› Process address space:

- Stack: temporary data extensible towards lower virtual addresses
- Heap: memory allocated dynamically extensible to higher virtual addresses
- Data section: global variable
- Text region: program code



## Process creation

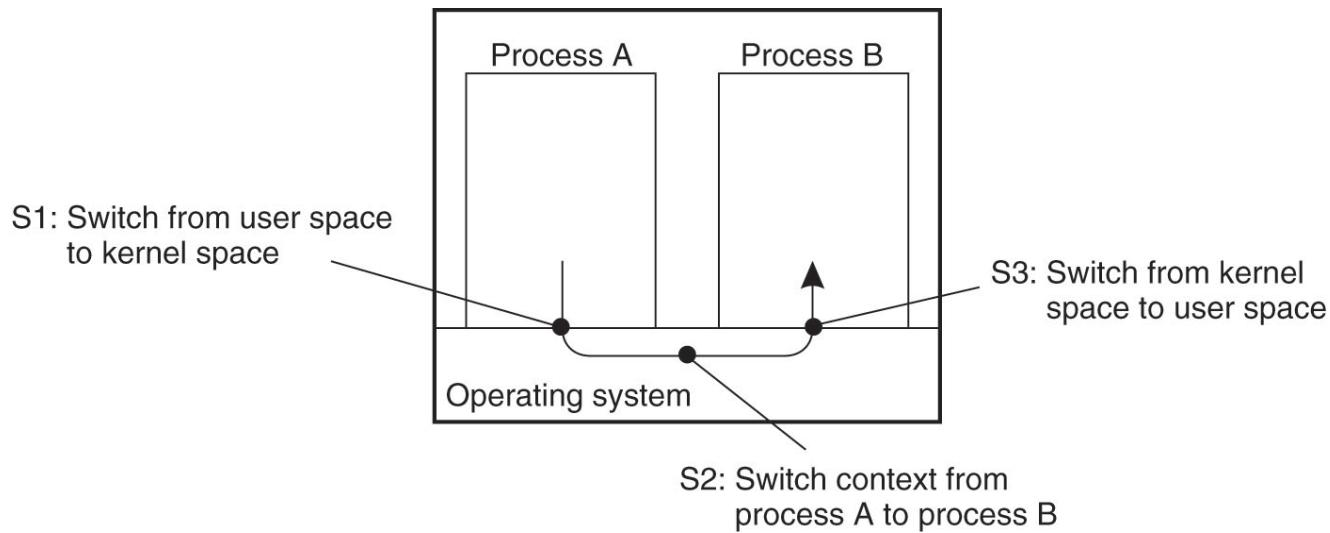
- › Process allocation: the act of choosing the host of the process
  - The System V (“5”) provides command to execute a process at an idle workstation
  - The Amoeba system chooses a host from a pool of processors to execute it
- › Execution environment: an address space w/ initialized content & open files
  - Static: program text, heap and stack regions created from a list
  - Dynamic: UNIX fork shares program text and copies stack and heap regions

## Interprocess communication (IPC):

- › Pipes/Message queues/Shared memory segments
- › Requires costly context switches

An interrupt requires kernel intervention

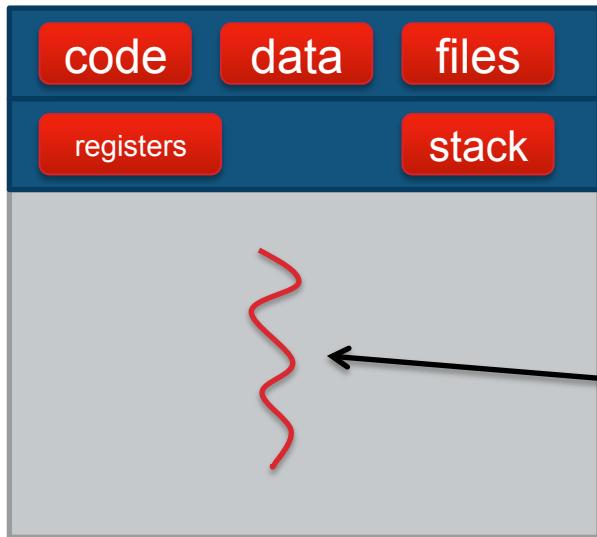
- › S1. Process A switches from user to kernel mode
  - Changing the memory map in the MMU (memory management unit)
  - Flushing the TLB (translation lookaside buffer)
- › S2. Context-switch (swapping processes from disk to main memory)
- › S3. Process B switches from kernel to user mode



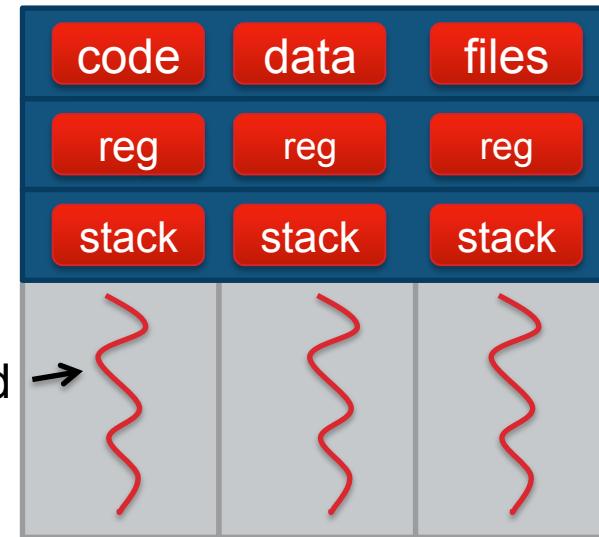
Threads: smallest unit of CPU utilization

› Portable OS Interface (POSIX) Threads

- Current activity: program counter
- Stack: temporary data
- No data, code, files (but it shares them with other threads)



Single-threaded process



Multithreaded process

*Threads:* smallest unit of CPU utilization

› POSIX Threads:

- Current activity: program counter
- Stack: temporary data
- No data, code, files (but it shares them with other threads)

› Communication between threads always through memory

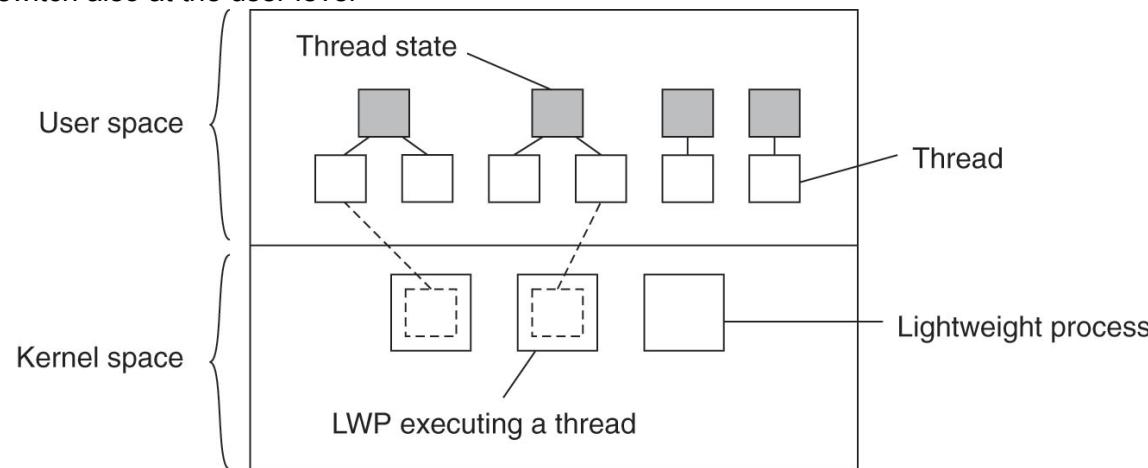
- Does not need IPC
- No context switches required (when purely at user-level)

## User-level threads library vs. kernel-scheduled thread

- › User-level thread library:
  1. Cheap to create and destroy
  2. Blocking system call freezes the entire process of the thread (and other threads)
  3. Cheap context-switch:
    - Few instructions to store and reload CPU register values
    - No need to change memory map in MMU
    - No need to flush the TLB
- › Kernel-scheduled threads:
  1. Costly to create and destroy
  2. Do not block the current process (and other threads) upon blocking system calls (I/O)
  3. Costly context-switch (similar to process context switch)
    - Needs to change memory map in MMU
    - Needs to flush the TLB

## User-level threads library vs. kernel-scheduled thread

- › Lightweight processes (LWP):
  1. Runs in the context of a process (potentially multiple LWPs per process)
  2. Current threads are in a table
  3. Once an LWP finds a runnable thread:
    - It locks the table in **user-mode** to update it
    - It switches context to that thread in **user-mode**
  4. **Does not block** upon system calls:
    - Passing from user to kernel mode in the LWP context
    - If the LWP is blocked, the kernel context-switches to another LWP
    - Implying a context-switch also at the user level





## › Processes

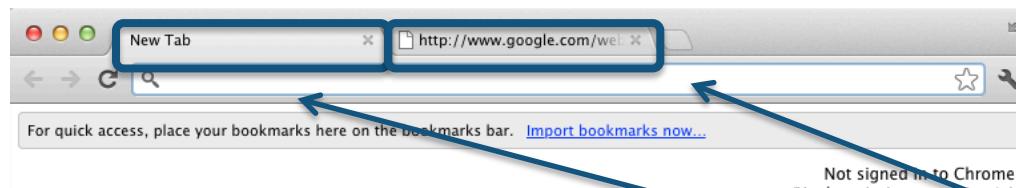
- **Isolated**: prevents one process from interfering with another
- **Inefficient**: starting/terminating a process and context switches are costly

## › Threads

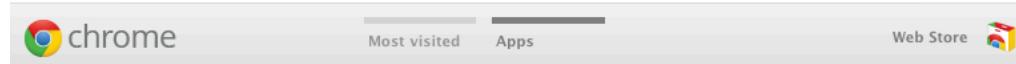
- **Non-isolated**: avoiding incorrect interferences makes programming harder
- **Efficient**: a thread is a lightweight version of a process

## Example: The Google Chrome web browser

- › Multithreading: each tab runs its own thread
- › A JavaScript error does not crash the main Chrome process:  
Only one tab freezes, the user can close it independently of others
- › Takes benefit of multicores architectures, each thread runs on a separate core



one distinct thread  
running each tab





THE UNIVERSITY OF  
SYDNEY

# Java Threads

They are more threads than you see

- › Java Virtual Machine (JVM) thread: main thread, responsible for executing VM operations
- › Periodic task thread: simulate timer interrupts to execute periodic operations, like scheduling other threads
- › Garbage collection threads: parallel and concurrent garbage collection
- › Compiler threads: runtime compilation of bytecode to native code
- › Signal dispatcher thread: redirects process directed signals to Java level signal handling methods
- › Graphical user interface: Java AWT and Swing creates threads for refreshing visual representation
- › Java Servlets and RMI: create pools of threads and invoke methods within these

## Synchronization calls

- › **thread.join([long milliseconds])**

Blocks the calling thread until the specified thread has terminated (for up to the specified milliseconds)

- › **thread.interrupt()**

Interrupts thread: causes it to return from a blocking method call such as sleep()

- › **object.wait([long milliseconds])**

Blocks the calling thread until a call made to notify() or notifyAll() on objects wakes the thread or the thread is interrupted (or the specified milliseconds have elapsed)

- › **object.notify()/object.notifyAll()**

Wakes, respectively, one (arbitrarily chosen) or all threads that have called wait on object.

## Thread lifecycle

### › NEW

A thread that is just instantiated is in *new* state. When a start() method is invoked, the thread moves to the ready state from which it is automatically moved to runnable state by the thread scheduler.

### › RUNNABLE (ready → running)

A thread executing in the JVM is ready to be placed in the running state by the scheduler

### › BLOCKED

A thread that is blocked waiting for a monitor lock is in this state. This can also occur when a thread performs an I/O operation and moves to next (runnable) state.

### › WAITING

A thread that is waiting indefinitely for another thread to perform a particular action is in this state.

### › SLEEPING

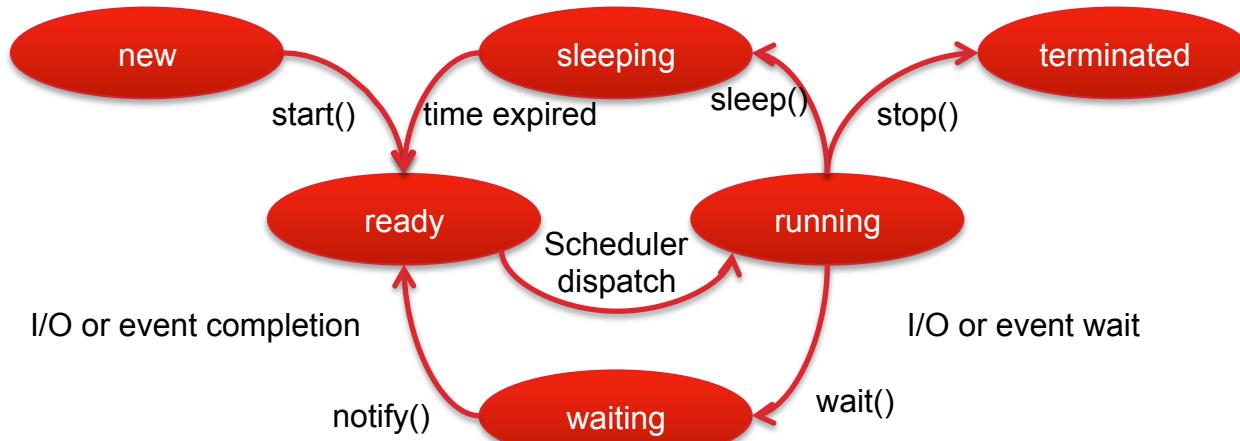
A thread that is waiting for another thread to perform an action for up to a specified waiting time is in this state.

### › TERMINATED (dead)

A thread that has exited is in this state.

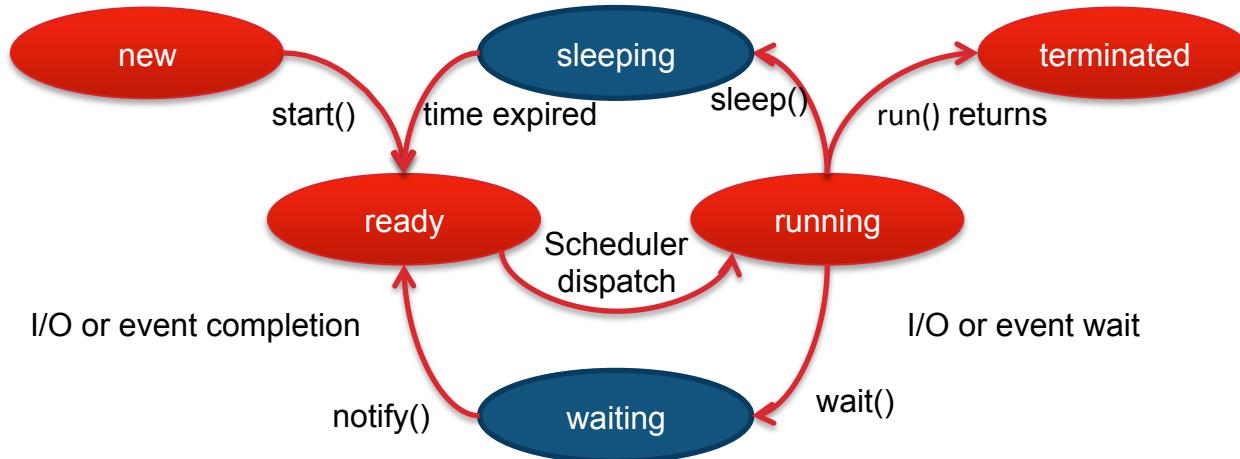


## Thread lifecycle (con't)





## Thread lifecycle (con't)





## Creating Java threads

There are two ways:

1. Deriving a class from the Thread class and overriding run()
  
2. Writing a class that implements the Runnable interface
  - Must redefine the run method
  - The code within the run method executes as a separate thread

```
class MyThread implements Runnable {  
    public void run() {  
        /* here goes my thread code */  
    }  
}
```



## Creating Java threads

There are two ways:

1. Deriving a class from the Thread class and overriding run()
  
2. Writing a class that implements the Runnable interface
  - Must redefine the run method
  - The code within the run method executes as a separate thread

```
class MyThread implements Runnable {  
    public void run() {  
        /* here goes my thread code */  
    }  
}
```



## Running Java threads

- › Instantiating a thread **is not sufficient** to run it

```
Thread mt = new Thread(new MyThread());  
mt.start();
```

- › One has to call its **start()** method, that will implicitly call its **run()** method



## Running Java threads

- › Instantiating a thread **is not sufficient** to run it

```
Thread mt = new Thread(new MyThread());  
mt.start();
```

- › One has to call its **start()** method, that will implicitly call its **run()** method

*Actually, it may not even start right after: but it is ready to be run by the scheduler*



- › A thread is generally faster than a process, but shares its data with other threads
- › To prepare a Java thread for execution, (1) implements the Runnable interface, (2) allocate a new thread, and (3) starts it!

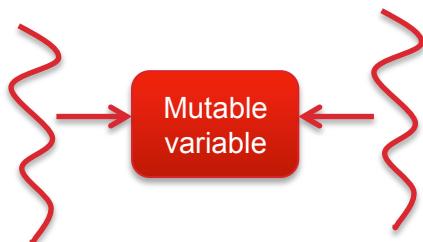


THE UNIVERSITY OF  
SYDNEY

# Multithreading

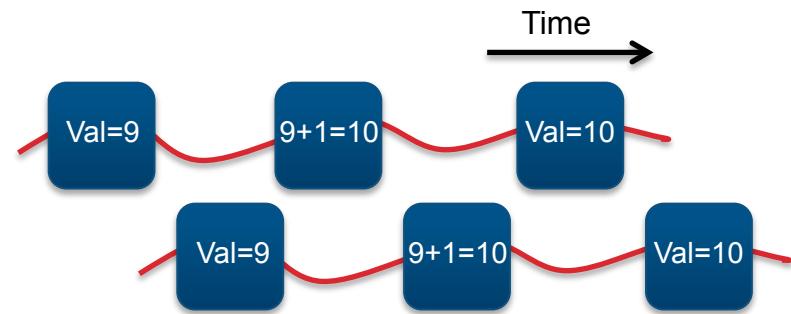
*Thread safety* is about protecting data from uncontrolled concurrent access

- › An *object state* is its data, stored in *state variables* such as instances or static fields
- › A *shared variable* is a variable that could be accessed by multiple threads
- › A *mutable variable* is a variable whose value may change
- › Whenever more than one thread access a given state variable and one of them might write to it, they all must coordinate their access to it using **synchronization**
- › If multiple threads access the same mutable state variable without appropriate synchronization, **the program is broken.**



```
public class UnsafeSequence {
    private int value;

    // Returns a unique value
    public int getNext() {
        return value++;
    }
}
```



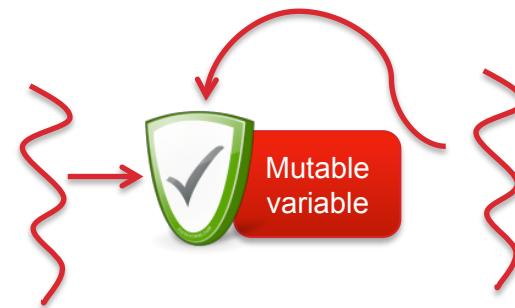
Multiple threads accessing a mutable variable may produce erroneous results

*Thread safety* is about protecting data from uncontrolled concurrent access

› Potential solutions:

- Don't share the state variable among threads, make it **thread specific**;
- Make the state variable **immutable**; or
- Use **synchronization** to access the state variable

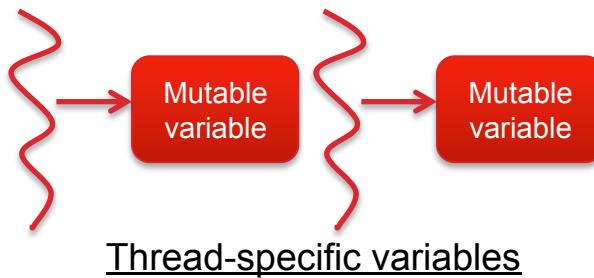
```
public class UnsafeSequence {  
    private int value;  
  
    // Returns a unique value  
    public synchronized int getNext() {  
        return value++;  
    }  
}
```



Using synchronization

## Thread-specific data

- › Threads belonging to the same process share data
- › How to declare data to be “thread-specific”?
  - By inheriting from Thread into MyThread in which data is declared
  - By using the ThreadLocal class initialized by initialValue()/set() and accessed through get()/set()

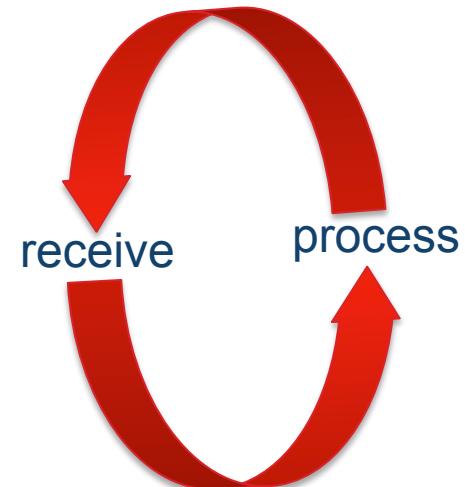


```
public class Service {  
    private static ThreadLocal errorCode = new ThreadLocal();  
  
    private static void transaction() {  
        try {  
            // some operation where an error may occur  
        } catch (Exception e) {  
            errorCode.set(e);  
        }  
  
        // get the error code for this transaction  
        public static Object getErrorCode() {  
            return errorCode.get();  
        }  
    }  
}
```

## Example: producer/consumer threads

### › Problem:

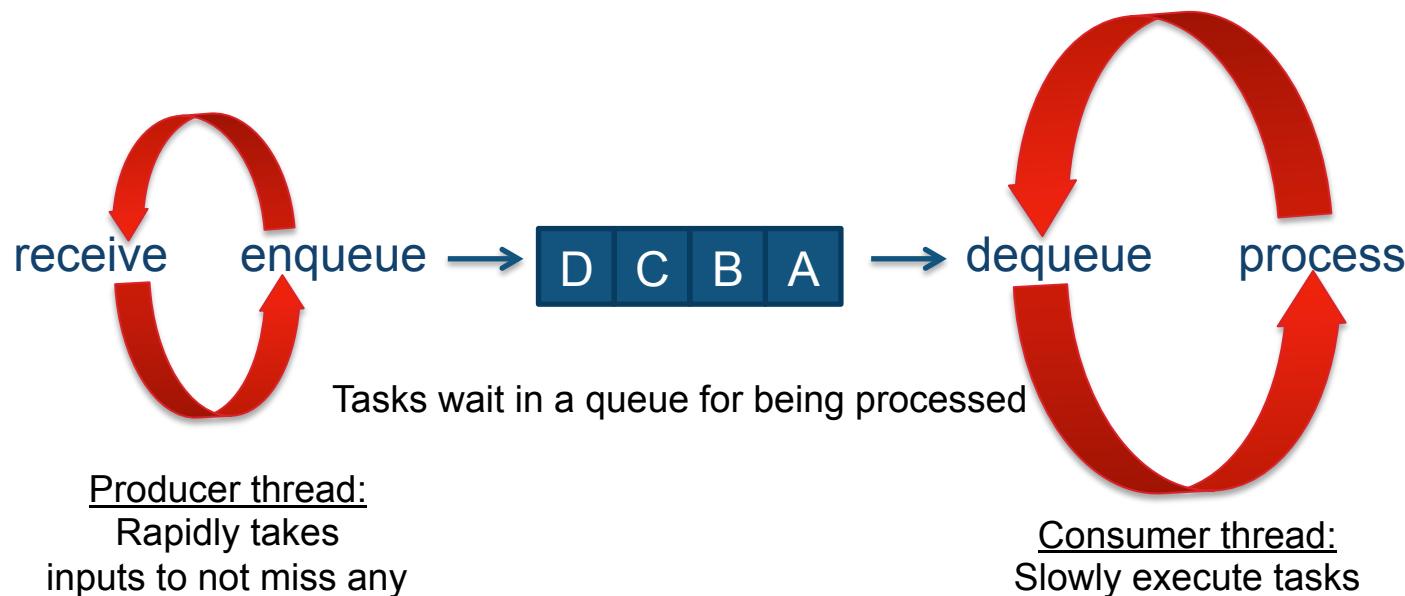
- Assume that you may receive **requests almost simultaneously**
- It takes you some **non-null amount of time** to process a request
- Yet, you would like to eventually **process all requests**
- How to proceed when one thread **cannot receive while processing** a request?



By the time, it can receive again,  
some request may be lost

### Example: producer/consumer threads (con't)

- › Solution: use multithreading in a consumer producer scenario
  - Ideal for **asynchronous** communication (mailbox analogy ≠ phone analogy) between threads
  - **Decouples** the consuming and production that run at different rates
  - One thread typically **passes the work** that arrives to some other thread
  - The producer loop can receive a request **while** the producer loop is processing another



## Example: producer/consumer threads (con't)

- › One thread acts as a consumer
- › Another thread acts as a producer

```
import java.util.Date;

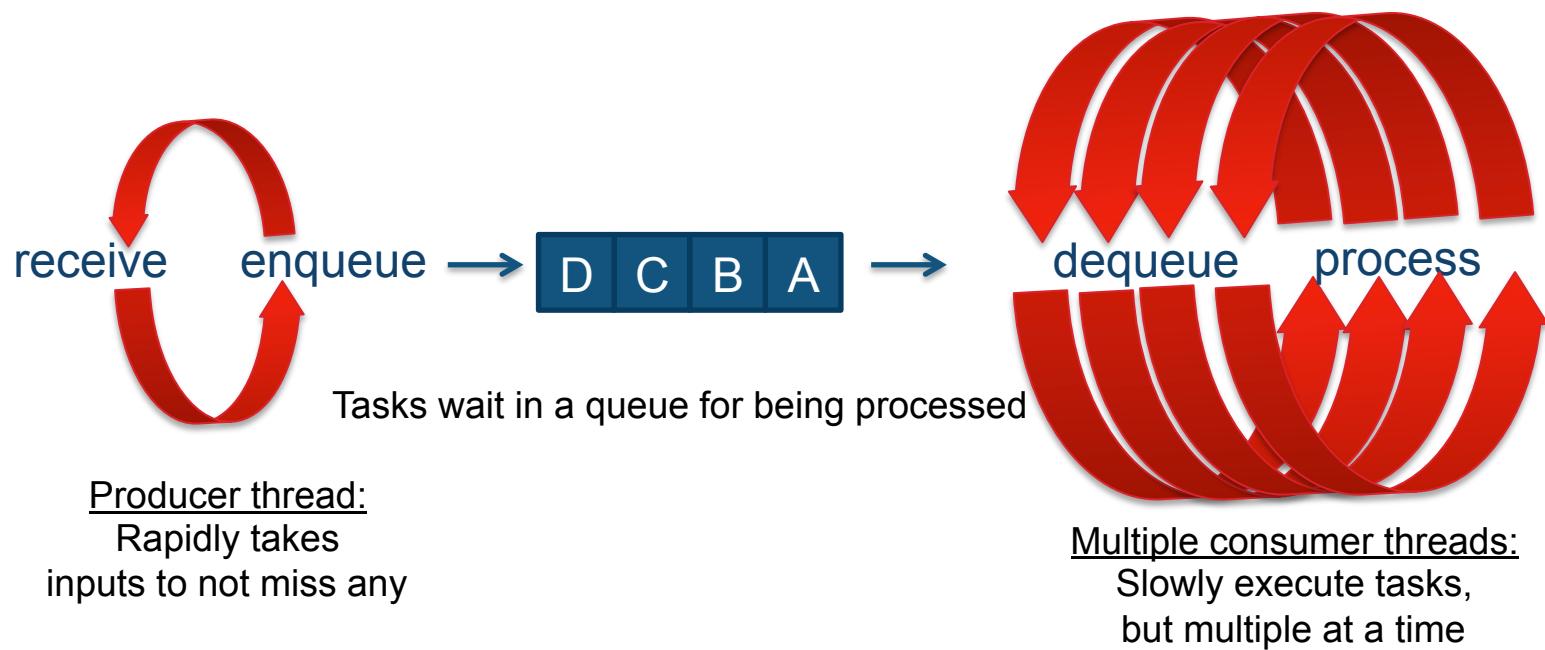
public class Factory {
    public static void main(String args[]) {
        // create the message queue
        Channel<Date> queue = new MessageQueue<Date>();

        // create the producer and consumer threads and pass
        // each thread a reference to the MessageQueue object.
        Thread producer = new Thread(new Producer(queue));
        Thread consumer = new Thread(new Consumer(queue));

        producer.start();
        consumer.start();
    }
}
```

## Example: producer/consumer threads (con't)

- › Speeding up the request processing using multiple consumer threads



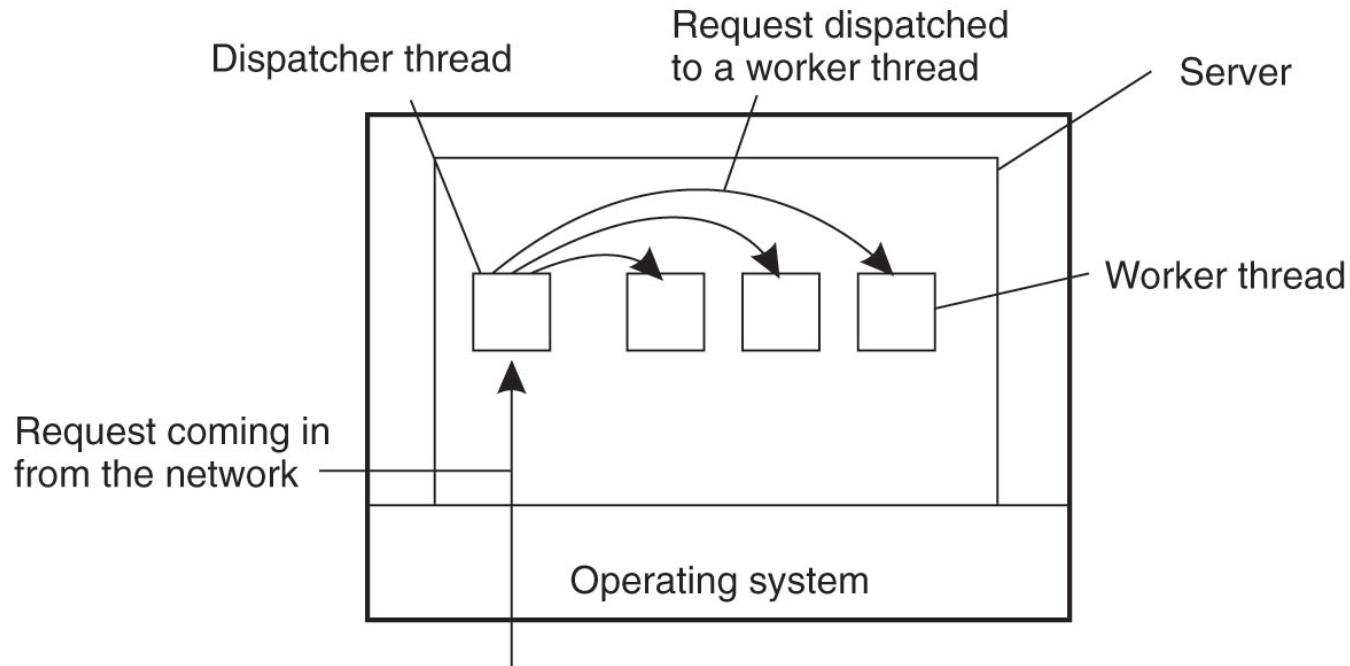


THE UNIVERSITY OF  
SYDNEY

# Multithreaded Server

A multithreaded server organized in a dispatcher/worker model

- › One thread, the *dispatcher*, reads incoming clients requests
- › The server chooses an *idle* worker thread and hands it the request



## Other alternatives

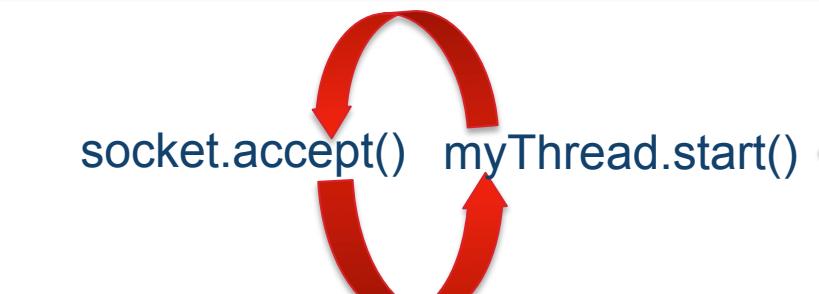
- › Single-thread server: one loop iteration must treat a request before starting to treat another one ⇒ potential request losses
- › Finite state machine: the server treats the request if the requested info is in the cache, otherwise it does not wait for I/O to disk to complete but stores the state of the request in a table. It uses non-blocking system calls.

Model	Characteristics
Threads	Parallelism, blocking system calls
Single-threaded process	No parallelism, blocking system calls
Finite-state machine	Parallelism, nonblocking system calls

## Creating threads upon request

- › For each request that arrives at the server
- › Creates the request in a separate thread

```
...  
    // keep treating incoming request  
    while (true) {  
        // block until accepting a new request  
        try { client.socket = this.serverSocket.accept();  
        } catch (IOException ioe) {System.out.println(ioe); }  
        // start a new task in a new thread  
        new Thread(new MyThread(new Task())).start();  
    }  
...
```



Main process thread: spawns one thread upon request



3 threads for 3 tasks



## Thread pool

- › Creates a pool of threads
- › That the server can (re)use
- › No need to initialize many threads
- › Nor to destroy many threads

```
import java.util.concurrent.*;  
  
public class TPExample {  
  
    public static void main(String[] args) {  
        int numTasks = Integer.parseInt(args[0].trim());  
  
        // create the thread pool  
        ExecutorService pool = Executors.newCachedThreadPool();  
  
        // Run each task using a thread in the pool  
        for (int i=0; i<numTasks; i++)  
            pool.execute(new Task());  
  
        // Shut down the pool. This shuts down the pool only  
        // after all threads have completed.  
        pool.shutdown();  
    }  
}
```

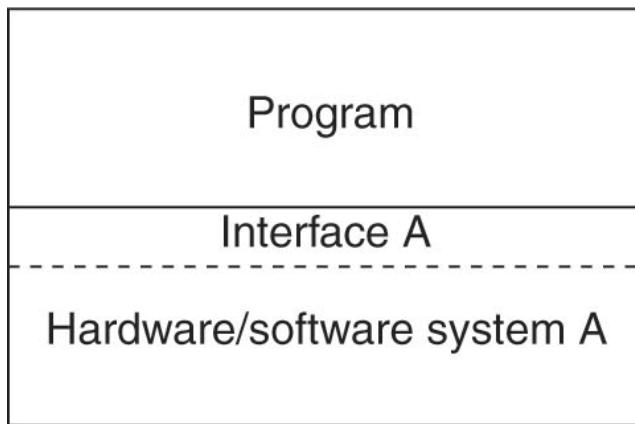


THE UNIVERSITY OF  
SYDNEY

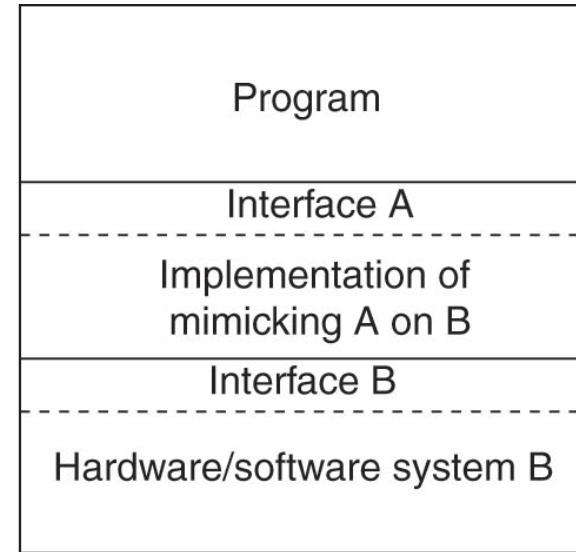
# Virtualization

*Virtualization:* changing an interface so as to mimic the behavior of a system

- › Every computer offers a programming interface to higher software
- › 1970's: multiple OSes should be able to run on a hardware (e.g., IBM 370 mainframe)
- › 1990's: multiple hardware should be able to run multiple software



(a)

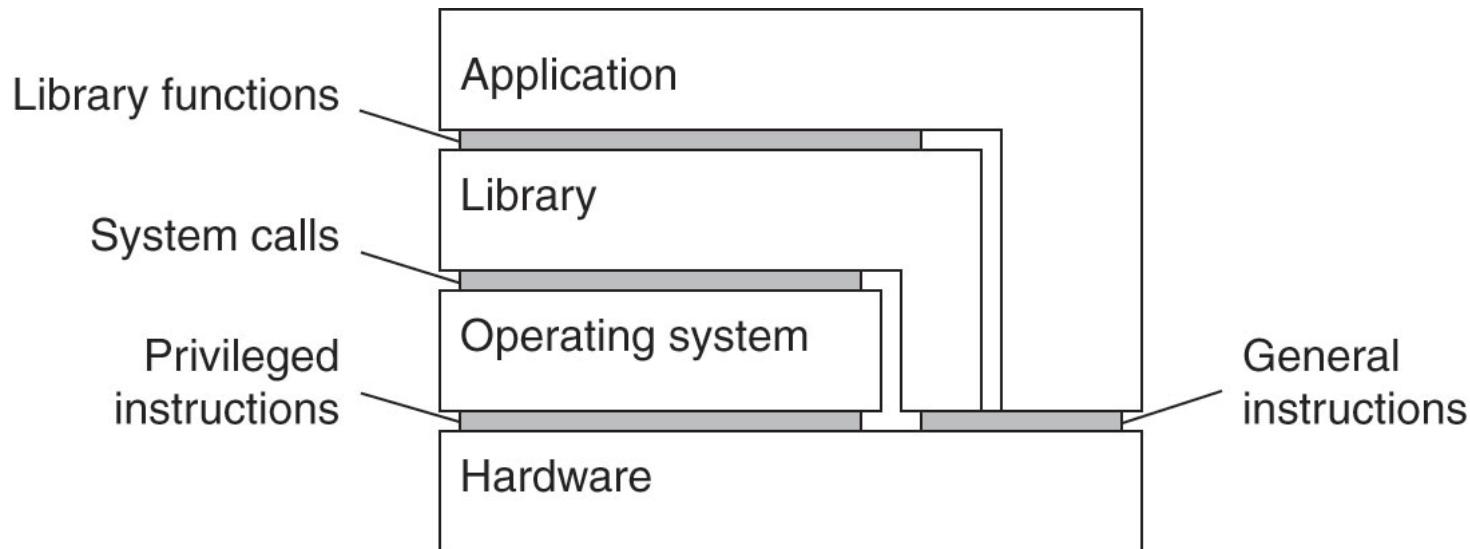


(b)

- › (a) General organization between a program an interface and a system
- › (b) General organization of virtualizing A on top of system B

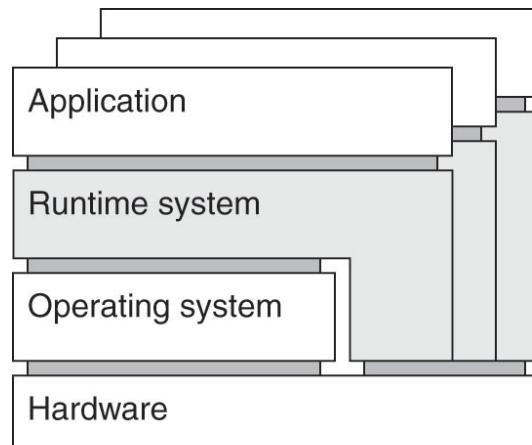
## Four types of interfaces at four different levels

- › Between hardware and software: machine instructions that can be invoked by any program
- › Between the hardware and software: machine instructions that can be invoked by privileged programs, like OS.
- › Between the OS and the library: system calls
- › Between library and application: library calls, generally forming what is known as an API

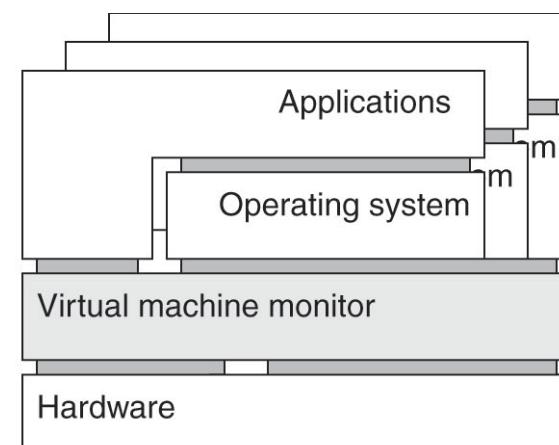


## Two types of virtual machines

- › Process virtual machine: A runtime system providing an abstract instruction set used by applications
  - Java Runtime Environment
  - Windows applications on UNIX platforms
- › Virtual Machine Monitor (VMM): provides a system implemented as a layer
  - VMware [Sugerman et al. 2001]
  - Xen [Barham et al. 2003]



(a)



(b)

- › (a) A process virtual machine, with multiple instances of (application, runtime) combinations
- › (b) A virtual machine monitor, with multiple instances of (application, operating system) combinations

## *Java Virtual Machine (JVM): specification of an abstract computer*

- › Java bytecode is not platform specific: can be run on top of almost any hardware
- › Javac produces some Java bytecode into .class files based on the .java source files
  
- › Class loader: loads the compiled .class files from both the Java program and the Java API
- › Java interpreter: executes the .class files by interpreting the bytecode and converting it into machine code to execute it
- › Just-in-time (JIT) compiler: during execution the just-in-time compiler identifies repetitive task to compile it into native machine instructions to speedup execution.
- › Java development kit (JDK): the mostly used software development kit for Java, comprises a set of resources to write Java program.
  
- › The Java API lists these features at <http://docs.oracle.com/javase/7/docs/api/>
- › In the HotSpot JVM, each Java thread is mapped to an OS thread, the OS maps OS threads to the underlying hardware resources (e.g., multi-cores)



- › Concurrency has been used for decades since hardware was powerful enough to address multiple humans need
- › Processes are heavy but protected whereas threads are lightweight but non-protected
- › Threads require synchronization to be protected from each other
- › Java uses a process virtual machine but maps its virtual Java threads to real OS threads